



education

phd | machine learning

uc berkeley | '17-'22

research: interpretable ml

advisor: bin yu

ms | data science

uc berkeley | '17-'21

research: statistical biology

bs | cs & math

university of virginia | '14-'17

double major

skills

machine learning

data cleaning • deep learning

pytorch • scikit-learn • jax

aws ec2 • s3 • sagemaker

keras • mllib • tensorflow

experienced

python • java • matlab

proficient

r • c/c++ • web basics

human languages

english • spanish • hindi

awards

berkeley grad slam semifinalist '19

outstanding teaching award '18

uva rader research award '17

uva undergrad symposium winner '17

raven honor society '16-'17

icpc regional qualification '14-'16

1st place microsoft code jam '16

3rd place google games uva '17

2nd place apt puzzle competition '17

funding awards

pdsoros fellowship finalist '19

ircn workshop travel award '19

vidya shelat fund award '16

rodman scholarship '14-'17

experience

paige ai | ai data scientist

summer '21

- interpretation and neural network modeling for medical imaging in oncology

berkeley | interpretable ml research (bin yu lab %)

fall '17 - present

- developed interpretation methods for ml models (e.g. neural nets)
- created interpretable models in medicine, biology, and computer vision

aws | research internship (pietro perona lab %)

summer '20

- testing for bias with causal matching using GANs
- interpreting semantic directions in generative models

response4life | volunteer data scientist

spring '20

- helped develop, integrate, and deploy models to forecast covid-19 severity

pacmed ai | interpretable ml internship

summer '19

- developed techniques to interpret machine-learning models for healthcare
- integrated interpretability techniques for predicting icu re-admission

facebook | computer vision internship

summer '17

- investigated unsupervised deep learning for segmentation of satellite imagery
- implemented crfs for segmentation post-processing

uva | ml research (yanjun qi lab %)

fall '16 – spring '17

- developed novel weighted- ℓ_1 , multi-task gaussian graphical model
- analyzed large-scale functional brain connectivity with graphical models

hhmi | ml research (srini turaga lab %)

summer '15, winter '15, summer '16

- extended cnns and watershed algorithms for neural image segmentation
- implemented distributed random forests for image segmentation

uva | comp. neuroscience research (william levy lab %)

fall '14 - fall '16

- developed detailed biophysical models of neural computation
- analyzed energy efficiency, noise, and variability in stochastic neurons

hhmi scientific computing | comp. neuroscience research

summer '14

- analyzed backpropagating action potentials via biophysical simulations

research innovations inc. | web dev + android internship

summer '13 - spring '14

- developed web/mobile app for task coordination with qr codes

coursework

computation

machine learning
computer vision
structure learning
algorithms
artificial intelligence
deep learning
learning theory
ai in graphics
cs theory
data structures
software dev. I & II
information retrieval
computer architecture

stat/math







statistical models
probability
statistics
optimization
linear algebra
info theory
real analysis
linear models
stochastic processes
chaos theory I & II
multivariate calculus
discrete mathematics
differential equations
abstract algebra

neuroscience





neural coding
neural network models
neurobiology
visual neuroscience
cognitive science

papers




interpretability

- adaptive wavelet distillation from neural networks through interpretations: ha, **cs**, et al. *neurips* '21 
- interpretations are useful: penalizing explanations to align neural networks with prior knowledge: rieger, **cs**, murdoch, & yu, *icml* '20 
- hierarchical interpretations for neural network predictions: **cs***, murdoch*, & yu, *iclr* '19 
- transformation importance with applications to cosmology: **cs***, ha*, lanusse, boehm, liu & yu, *iclr* '20 *workshop (spotlight talk)* 
- interpretable machine learning: definitions, methods, and applications: murdoch*, **cs***, kumbier, abbasi-asl, & yu, *pnas* '19 
- disentangled attribution curves for interpreting random forests and boosted trees: devlin, **cs**, & yu *arxiv* '19 



interpretable data science projects

- curating a covid-19 data repository and forecasting county-level death counts in the united states: altieri, barter, ..., **cs***, ..., & yu* *harvard data science review* '20  
- developing reliable clinical decision rules: a case study in identifying blunt abdominal trauma in children: kornblith*, **cs***, et al. *seam abstract* 
- interpretable deep learning for accurate molecular partner prediction in clathrin-mediated endocytosis: **cs***, li* et al. *in prep* 





software packages

- imodels: a python package for interpretable modeling: **cs***, nasser*, tan, tang, & yu, *journal of open source software* '21  
- veridical-flow: a python package for facilitating stable data analysis 



causal inference / ml theory

- matched sample selection with GANs for mitigating attribute confounding: **cs**, balakrishnan, & perona *cvpr* '21 *civ workshop* 
- revisiting complexity and the bias-variance tradeoff: dwivedi*, **cs***, yu, & wainwright *topml workshop* '21 

statistical neuroscience


- large scale image segmentation with structured-loss-based deep learning for connectome reconstruction: funke et al. *tpami* '18 
- a weighted- ℓ_1 , multi-task graphical model with applications to heterogeneous brain connectivity: **cs**, wang, & qi, *neurips* '17 *amlicd workshop* 
- a consensus layer V pyramidal neuron can sustain interpulse-interval coding: **cs** & levy *plos one* '17 
- linearized synaptic integration at no extra cost: morel, **cs**, & levy, *journal of computational neuroscience* '18 

selected talks

- interpreting ml models: *uc berkeley bair seminar*, '20 
- uncovering brain connections underlying autism via graphical models: *tom tom founder's machine learning conference*, '17 

teaching





berkeley | student instructor

summer 2018
machine learning: cs 189/289 
lectures to class of 80+ students

fall 2019

artificial intelligence: cs 188 

projects / activities

notes, blog, & slides  '14-'20
hummingbird tracking  '18
news balancer django app  '17
java mini-games  '14-'16

service

basis education volunteering '19-'20
bair undergrad mentoring '18-'20
neurips reviewer '21
iclr workshop reviewer '21
cvpr reviewer '21
aaai xai workshop reviewer '21
neurips ml4h workshop reviewer '20
computer literacy volunteering '15-'17